

Flash R

Minerva Statistical Consulting

David John Baker, Ph.D.

07/20/2019



Session I (< 45 Minutes)

Afternoon Goals

- Install R
- Install RStudio
- Run Basic Commands in Console
- Run Basic Tidyverse Commands

Why R?

1. The R community is fantastic, check out #rstats on Twitter as well as everyone affiliated with the tidyVerse
2. R will always be free because the people behind it believe in open source principles.
3. Time spent learning R is time spent learning how computers work. If you learn about R, you are also learning computer programming. Time spent in something like SPSS or SAS does not easily transfer to other programs.
4. On r-jobs.com the way they decide to split jobs is jobs that make above and below \$100,000.
5. R is your ticket out of academia, if you need it. It's also insane to think people would learn so much about statistics, the hardest part about becoming a data scientist, without learning the software to get you in the door.
6. When you make analyses and graphs in R they are very easy to reproduce. You just press 'Run' again.
7. If you do your data cleaning in R, then each step is documented. There is less chance for human error.
8. It makes gorgeous graphs.
9. There are a lot of ways that R integrates into other software. This book is written in bookdown, my website is written in blogdown, you can also make interactive data applications.
10. It's kind of fun.



[CRAN](#)
[Mirrors](#)
[What's new?](#)
[Task Views](#)
[Search](#)

[About R](#)
[R Homepage](#)
[The R Journal](#)

[Software](#)
[R Sources](#)
[R Binaries](#)
[Packages](#)
[Other](#)

[Documentation](#)
[Manuals](#)
[FAQs](#)
[Contributed](#)

The Comprehensive R Archive Network

Download and Install R

Precompiled binary distributions of the base system and contributed packages, **Windows and Mac** users most likely want one of these versions of R:

- [Download R for Linux](#)
- [Download R for \(Mac\) OS X](#)
- [Download R for Windows](#)

R is part of many Linux distributions, you should check with your Linux package management system in addition to the link above.

Source Code for all Platforms

Windows and Mac users most likely want to download the precompiled binaries listed in the upper box, not the source code. The sources have to be compiled before you can use them. If you do not know what this means, you probably do not want to do it!

- The latest release (2018-12-20, Eggshell Igloo) [R-3.5.2.tar.gz](#), read [what's new](#) in the latest version.
- Sources of [R alpha and beta releases](#) (daily snapshots, created only in time periods before a planned release).
- Daily snapshots of current patched and development versions are [available here](#). Please read about [new features and bug fixes](#) before filing corresponding feature requests or bug reports.
- Source code of older versions of R is [available here](#).
- Contributed extension [packages](#)

Figure 1: CRAN Homepage

R, RStudio, and Tidyverse

R

You can download R for your computer by going to CRAN and selecting the appropriate **Download and Install R** links. Make sure to install R first before installing RStudio.

RStudio


RStudio is an integrated development environment (IDE) for R¹. RStudio is basically your workbench where you can access everything you need for managing your scripts, data, and project structure. By using RStudio, you also can use a host of other features ranging from Markdown documents (like this one!), interactive data dashboards like Shiny, and the tidyverse.


RStudio Environment

Once you now have R and RStudio installed, it's time to open up RStudio. By opening RStudio, you are also starting R. R will be running under the hood of RStudio. After installing R, you can run it on it's own by typing **R** into your terminal on a Unix machine (Mac, Linux). Though after seeing how RStudio works, you would realize why doing this is basically masochistic. (If you do this, you can quit out of the terminal R with `quit()` followed by `n`).

When you first open RStudio will see a few different panels. In it's default settings, the bottom left is the Console. The top right has your Environment, History, and version control commands. The bottom right has your Viewer, Library for your packages, and a system to navigate your files. The top left will be where you write your code.


¹<https://www.rstudio.com/products/RStudio/>



[Products](#)
[Resources](#)
[Pricing](#)
[About Us](#)
[Blogs](#)


Choose Your Version of RStudio

RStudio is a set of integrated tools designed to help you be more productive with R. It includes a console, syntax-highlighting editor that supports direct code execution, and a variety of robust tools for plotting, viewing history, debugging and managing your workspace. [Learn More about RStudio features.](#)



RStudio Desktop Open Source License	RStudio Desktop Commercial License	RStudio Server Open Source License	RStudio Server Pro Commercial License	RStudio Server Pro + RStudio Connect Commercial License
FREE	\$995 per year	FREE	\$9,995 per year	\$29,995 per year
DOWNLOAD Learn More	BUY Learn More	DOWNLOAD Learn More	DOWNLOAD Learn More	TALK Learn More

Figure 2: RStudio

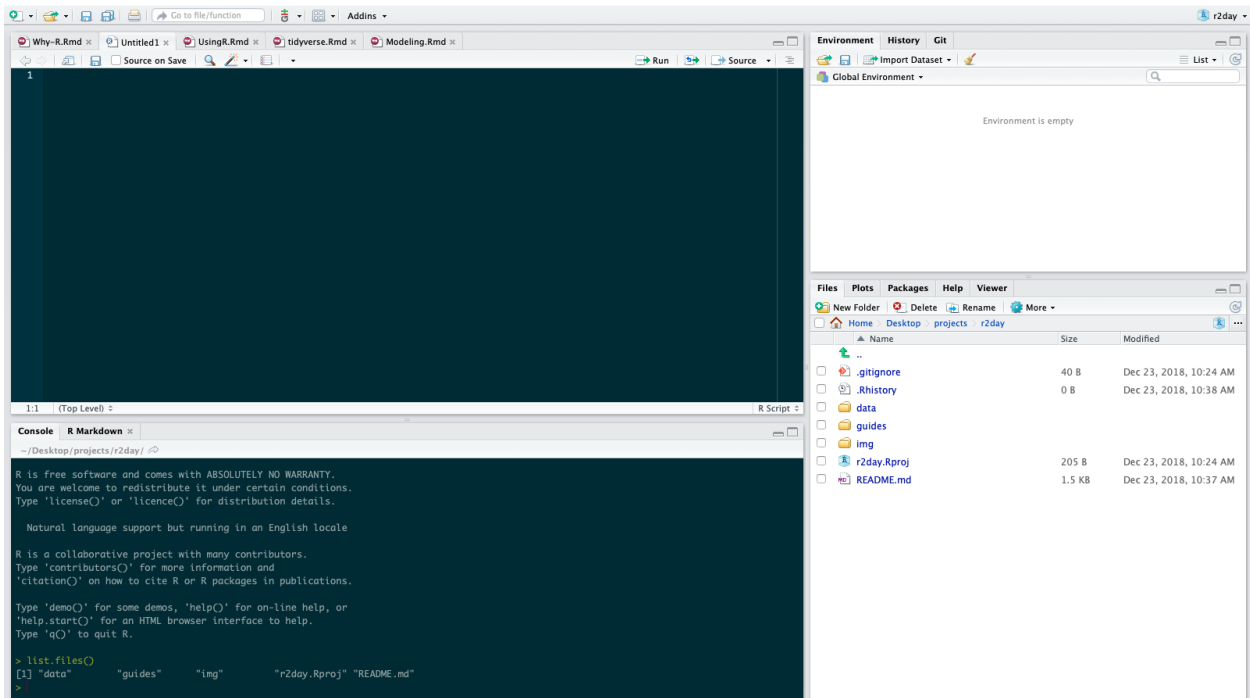


Figure 3: RStudio Environment

Environment

The top left has information about your current Environment. As you make new things in an R session you can track them here. There is also a History tab here that keeps track of code you wrote. Additionally there is a Git tab that will eventually allow you to do version control. You don't have to know what that is, but one day you might read about it.

Viewer

The bottom right is your File Explorer/Finder window. Try to click around on the **Files** tab. When you click **Plots** there should be nothing there as you have not made any plots yet. Your **Packages** tab will have a listing of software that you can load into R. Notice that if you click one of the package names, it will navigate you to the **Help** tab. Lastly, the Viewer tab will let you display any documents that you make while writing in R. This could be markdown documents or maybe a website that you are writing eventually.

It is important to note that you will probably “break” R and RStudio many times when learning. Know that this is OK and the some of the best advice for learning how to program is by just seeing what happens when you change something and Googling your problems.

Console

The Console in R is where you can run one-off R commands. Try to type a few of the following commands into the Console.

```
list.files()
```

```
## [1] "Flash_R_files" "Flash_R.html" "Flash_R.Rmd" "flash_R.Rproj"
## [5] "img"           "README.md"     "slides"
```

```
str(iris)
```

```
## 'data.frame': 150 obs. of 5 variables:
## $ Sepal.Length: num 5.1 4.9 4.7 4.6 5 5.4 4.6 5 4.4 4.9 ...
## $ Sepal.Width : num 3.5 3 3.2 3.1 3.6 3.9 3.4 3.4 2.9 3.1 ...
## $ Petal.Length: num 1.4 1.4 1.3 1.5 1.4 1.7 1.4 1.5 1.4 1.5 ...
## $ Petal.Width : num 0.2 0.2 0.2 0.2 0.2 0.4 0.3 0.2 0.2 0.1 ...
## $ Species : Factor w/ 3 levels "setosa","versicolor",...: 1 1 1 1 1 1 1 1 1 1 ...
```

```
2 + 2
```

```
## [1] 4
```

Each of these will create a different kind of output. Now try to put something in your R console that will create an error message. Maybe some math that ends with an operation sign? Maybe some text? In the next session, we will go over what is legal and illegal input in the R Console.

Editor

The top left panel is where you edit your documents. RStudio allows you to handle many different types of documents. In this course we will mostly use RMarkdown files. These files end in `.Rmd` and allow for both text and R code. R scripts on the other hand only handle R code.

Using the editor, you should also familiarize yourself with the keyboard shortcuts in RStudio. For example, to run a line of code in the Editor, you can press **CMD + ENTER** on Mac or **CTRL + ENTER** on Windows. When the cursor is on a line that has runnable R code, this will run that line in the console. You can also use your

mouse highlight many lines of R code and run the same commands. We will get a lot of practice with this in the next session.

Session 2 - Using R (1 Hour)



Lesson Goals

- Run Basic Commands in R
- Understand Basic Data Structures
- Run commands over vectors
- Index data frames
- Learn basic data structures
- Understand base R vs Tidyverse
- Import and export data to/from R

RStudio Shortcuts and Markdown

In this session I will be using a lot of Keyboard Shortcuts when typing myself. In the past, people have always asked about these, so I'm anticipating that question with a link here to that page.

R as Calculator

The Console of R is where all the action happens. You can use it just like you would use a calculator. Try to do some basic math operations in it.

```
2 + 2
```

```
## [1] 4
```

```
5 - 2
```

```
## [1] 3
```

```
10 / 4
```

```
## [1] 2.5
```

```
9 * 200
```

```
## [1] 1800
```

```
sqrt(81)
```

```
## [1] 9
```

```
10 > 4
```

```
## [1] TRUE
```

```
2 < -1
```

```
## [1] FALSE
```

Now from the output above, you'll notice that there are a few different types of responses that R will give. For the math responses, we get numbers, but we can also get TRUE and FALSE statements.

When working with data, we need to be aware not only of what the data represents, but what R thinks it represents. We won't go over the differences between things like ordinal, ratio, and categorical data, I'll assume you have a basic understanding of this. What we will focus on is the different data types that R thinks in.

For now, we are going to talk about R's basic data structures.

- Logical
- Integer
- Double (numeric)
- Character
- Factor

The first is logical. Logical is basically just TRUE or FALSE. We can try a few different expressions that show how this works.

```
2 > 4
```

```
## [1] FALSE
```

```
1 > 0
```

```
## [1] TRUE
```

```
4 >= 7
```

```
## [1] FALSE
```

```
5 != 5
```

```
## [1] FALSE
```

Eventually you will learn to take advantage of the complexities of this when we get to subsetting and combining them with other logical operators like &(and) and | (OR).

Next we have integers and double. Both integers and double are R's numeric forms of data. The `is.numeric()` command checks for if data is number-y.

```
is.integer(7L)
```

```
## [1] TRUE
```

```
is.double(7)
```

```
## [1] TRUE
```

```
is.numeric(7)
```

```
## [1] TRUE
```

Next we have characters. Characters are not *just* letters, but rather data that is text. Character data is always wrapped in quotes " "

```
is.character("hello, world!")
```

```
## [1] TRUE
```

```
is.character("7")
```

```
## [1] TRUE
```

```
is.character("I will drink 7 coffees by the end of today!")
```

```
## [1] TRUE
```

```
is.character("NA")
```

```
## [1] TRUE
```

Note that if a special character like NA is in quotes, R will still think it is a character. To change this, we need to coerce our data into a different type. We will cross that bridge later. For now, you just need to be aware of the different character types.

Lastly, there are factors which sometimes LOOK like characters, but are R's way of thinking about categorical data. We need to assign this to R. When you first import in data into R, it will sometimes guess it as being a factor which is very annoying! If R is being slow, or not responding to something you want it to do, a common rookie mistake is to have your data accidentally be a factor.

```
is.factor("doggo")
```

```
## [1] FALSE
```

```
doggo <- as.factor("doggo")
```

```
is.factor(doggo)
```

```
## [1] TRUE
```

```
is.character(doggo)
```

```
## [1] FALSE
```

```
is.numeric(doggo)
```

```
## [1] FALSE
```

Now that we're at least aware of the different types of data in R, we can move on to building up an intuitive understanding of how R thinks about data under the hood.

Being Lazy

You don't always want to print your output and retype it in. The idea of being a good programmer is to be very lazy (efficient).

One of the best ways to be efficient when programming is to save variables to objects. Below is some example code that uses the <- operator to assign some math to an object. After you assign it to an object, you can then manipulate it like you would any other number. Yes, you can use = as an assignment operator (for all you Pythonistas), but in R this is considered bad practice as R is primarily a statistical programming language and the = sign means something very different in a math context.

```
foo <- 2 * 3
```

```
foo * 6
```

```
## [1] 36
```

After running these two lines of code, notice what has popped up in your environment in RStudio! You should see that you now have an object in the Environment called `foo`.

In addition to saving single values to objects, you can also store a collection of values. Below we use an example that might have a bit more meaning, the below stores what could be some data into an object that represents what it might be.

```
yearsSellingWidgets <- c(2,1,4,5,6,7,3,2,4,5,3)
```

The way that the line above works is that we use the `c()` function (c for combine) to group together a bunch of the same type of data (numbers) into a vector. Once we have everything combined and stored into an object, we can then manipulate all the numbers in the object just like we did above with a single number. A single dimensional object is called a **vector**. For example, we could multiply all the numbers by three.

```
yearsSellingWidgets * 3
```

```
## [1] 6 3 12 15 18 21 9 6 12 15 9
```

Or maybe we realized that our inputs were wrong and we need to shave off two years off of each of the entries.

```
yearsSellingWidgets - 2
```

```
## [1] 0 -1 2 3 4 5 1 0 2 3 1
```

Or perhaps we want to find out which of our pieces of data (and other data associated with that observation) are less than 2.

```
yearsSellingWidgets < 2
```

```
## [1] FALSE TRUE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
```

Any sort of mathematical operation can be performed on a vector! In addition to treating it like a mathematical operation, we can also run functions on objects. By looking at the name of each function and its output, take a guess at what each of the below functions does.

```
mean(yearsSellingWidgets)
```


```
## [1] 3.818182
```

```
sd(yearsSellingWidgets)
```

```
## [1] 1.834022
```

```
hist(yearsSellingWidgets)
```


Histogram of yearsSellingWidgets



```
scale(yearsSellingWidgets)
```

```
##           [,1]
## [1,] -0.99136319
## [2,] -1.53661295
## [3,]  0.09913632
## [4,]  0.64438608
## [5,]  1.18963583
## [6,]  1.73488559
## [7,] -0.44611344
## [8,] -0.99136319
## [9,]  0.09913632
## [10,] 0.64438608
## [11,] -0.44611344
## attr("scaled:center")
## [1] 3.818182
## attr("scaled:scale")
## [1] 1.834022
```

```
range(yearsSellingWidgets)
```

```
## [1] 1 7
```

```
min(yearsSellingWidgets)
```

```
## [1] 1
```

```
class(yearsSellingWidgets)
```

```
## [1] "numeric"
```

```
str(yearsSellingWidgets)
```

```
##  num [1:11] 2 1 4 5 6 7 3 2 4 5 ...
```

```
summary(yearsSellingWidgets)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      1.000   2.500   4.000   3.818   5.000   7.000
```

Often working with data, we don't want to just play with one group of numbers. Most of the time we are trying to compare different observations in data science. If we then create two vectors (one of which we have already made!) and then combine them together into a data frame, we have something sort of looking like a spreadsheet. A two-dimensional object is called a **data frame**.

```
yearsSellingWidgets <- c(2,1,4,5,6,7,3,2,4,5,3)
numberOfSales <- c(5,2,5,7,9,9,2,8,4,7,2)
salesData <- data.frame(yearsSellingWidgets,numberOfSales)
salesData
```

```
##      yearsSellingWidgets numberOfSales
## 1                      2             5
## 2                      1             2
## 3                      4             5
## 4                      5             7
## 5                      6             9
## 6                      7             9
## 7                      3             2
## 8                      2             8
## 9                      4             4
## 10                     5             7
## 11                     3             2
```

Now if we wanted to use something like R's correlation function we could just pass in the two objects that we have like this and get a correlation value.

```
cor(yearsSellingWidgets,numberOfSales)
```

```
## [1] 0.6763509
```

But often our data will be saved in data frames and we need to be able to access one of our vectors inside our data frame. To access a piece of information in a data frame we use the `$` operator.

```
salesData$yearsSellingWidgets
```

```
## [1] 2 1 4 5 6 7 3 2 4 5 3
```

Running the above code will print out the vector called `yearsSellingWidgets` from the data frame `salesData`. Using this form, we can then use this with the correlation function.

```
cor(salesData$yearsSellingWidgets,salesData$numberOfSales)
```

```
## [1] 0.6763509
```


In addition to just getting numeric output, we also want to be able to look at our data. Take a look at the code below and try to figure out what the function call is, as well as what each argument (or thing you pass to a function) does.

```
plot(yearsSellingWidgets,numberOfSales,
     data = salesData,
     main = "My Plot",
```

```
xlab = "Years at Company",  
ylab = "Number of Sales")
```

```
## Warning in plot.window(...): "data" is not a graphical parameter  
## Warning in plot.xy(xy, type, ...): "data" is not a graphical parameter  
## Warning in axis(side = side, at = at, labels = labels, ...): "data" is not  
## a graphical parameter  
  
## Warning in axis(side = side, at = at, labels = labels, ...): "data" is not  
## a graphical parameter  
  
## Warning in box(...): "data" is not a graphical parameter  
## Warning in title(...): "data" is not a graphical parameter
```

My Plot



If you are having a hard time understanding arguments, one thing that might help to think about is that each argument is like a click in a software program like SPSS or Excel. Imagine you want to make the same plot with this data in SSPSS, what would you do? The first thing you would do is to go to the top of the bar and find the **Plot** function and click it. This is the same as typing out `plot()` in R. Then you would have to tell that new pop up screen what two variables you want to plot and click on the related variables. Dragging and dropping those variables into your plot builder is the same as just typing out the variables you want. Lastly you want to put names on your axes and a title on your plot. The same logic would follow. We'll explore these ideas a bit more in the next section

Packages and Help

One of best things to do is just open an R help page and play around with things (and break things) until you “get” how it works.

How to actually learn any new programming concept



Essential

Changing Stuff and
Seeing What Happens

O RLY?

@ThePracticalDev

Figure 4: CRAN Homepage

To access R's in built help function you can easier use the Help viewer in R studio or type in a question mark before the command in the console. Using two ?? will search more generally

```
?scale()
??scale()
```

Data Exploration

```
library(tidyverse)
```

```
## -- Attaching packages ----- tidyverse 1.2.1 --

## v ggplot2 3.1.0      v purrr   0.3.2
## v tibble  2.1.1      v dplyr  0.8.0.1
## v tidyr   0.8.3      v stringr 1.4.0
## v readr   1.3.1      v forcats 0.3.0

## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
```

```
str(txhousing)
```

```
## Classes 'tbl_df', 'tbl' and 'data.frame':  8602 obs. of  9 variables:
## $ city      : chr  "Abilene" "Abilene" "Abilene" "Abilene" ...
## $ year      : int   2000 2000 2000 2000 2000 2000 2000 2000 2000 2000 ...
## $ month     : int    1  2  3  4  5  6  7  8  9 10 ...
## $ sales     : num   72  98 130  98 141 156 152 131 104 101 ...
## $ volume    : num  5380000 6505000 9285000 9730000 10590000 ...
## $ median    : num   71400 58700 58100 68600 67300 66900 73500 75000 64500 59300 ...
## $ listings  : num   701 746 784 785 794 780 742 765 771 764 ...
## $ inventory : num    6.3 6.6 6.8 6.9 6.8 6.6 6.2 6.4 6.5 6.6 ...
## $ date      : num   2000 2000 2000 2000 2000 ...
```

```
class(txhousing)
```

```
## [1] "tbl_df"      "tbl"        "data.frame"
```

```
summary(txhousing)
```

```
##      city      year      month      sales
## Length:8602    Min.   :2000    Min.   : 1.000    Min.   :  6.0
## Class :character 1st Qu.:2003    1st Qu.: 3.000    1st Qu.: 86.0
## Mode  :character Median :2007    Median : 6.000    Median : 169.0
##              Mean  :2007    Mean  : 6.406    Mean   : 549.6
##              3rd Qu.:2011    3rd Qu.: 9.000    3rd Qu.: 467.0
##              Max.   :2015    Max.   :12.000    Max.   :8945.0
##              NA's   :568
##      volume      median      listings      inventory
## Min.   :8.350e+05    Min.   : 50000    Min.   :  0    Min.   : 0.000
## 1st Qu.:1.084e+07    1st Qu.:100000    1st Qu.: 682    1st Qu.: 4.900
## Median :2.299e+07    Median :123800    Median : 1283    Median : 6.200
## Mean   :1.069e+08    Mean   :128131    Mean   : 3217    Mean   : 7.175
## 3rd Qu.:7.512e+07    3rd Qu.:150000    3rd Qu.: 2954    3rd Qu.: 8.150
## Max.   :2.568e+09    Max.   :304200    Max.   :43107    Max.   :55.900
## NA's   :568         NA's   :616      NA's   :1424    NA's   :1467
```

```
##      date
## Min.   :2000
## 1st Qu.:2004
## Median :2008
## Mean   :2008
## 3rd Qu.:2012
## Max.   :2016
##
```

```
# View(txhousing)
```


Accessing individual 'columns' is done with the \$ operator

```
txhousing$sales
```

Can you use this to plot the different numeric values against each other?

What would the follow commands do?

```
hist(scale(iris$Sepal.Length))
```



```
iris$Sepal.Length.scale <- scale(iris$Sepal.Length)
```

Indexing

Let's combine logical indexing with creating new objects.

What do the follow commands do? Why?

```
txhousing[1,1]
```

```
## # A tibble: 1 x 1
```

```
## city
## <chr>
## 1 Abilene
```

```
txhousing[2,]
```

```
## # A tibble: 1 x 9
##   city    year month sales  volume median listings inventory date
##   <chr>  <int> <int> <dbl>   <dbl>   <dbl>   <dbl>   <dbl> <dbl>
## 1 Abilene 2000     2    98 6505000  58700     746     6.6 2000.
```

```
txhousing[,5]
```

```
## # A tibble: 8,602 x 1
##   volume
##   <dbl>
## 1 5380000
## 2 6505000
## 3 9285000
## 4 9730000
## 5 10590000
## 6 13910000
## 7 12635000
## 8 10710000
## 9 7615000
## 10 7040000
## # ... with 8,592 more rows
```

```
txhousing[txhousing$year < 2003,]
```

```
## # A tibble: 1,656 x 9
##   city    year month sales  volume median listings inventory date
##   <chr>  <int> <int> <dbl>   <dbl>   <dbl>   <dbl>   <dbl> <dbl>
## 1 Abilene 2000     1    72 5380000  71400     701     6.3 2000
## 2 Abilene 2000     2    98 6505000  58700     746     6.6 2000.
## 3 Abilene 2000     3   130 9285000  58100     784     6.8 2000.
## 4 Abilene 2000     4    98 9730000  68600     785     6.9 2000.
## 5 Abilene 2000     5   141 10590000 67300     794     6.8 2000.
## 6 Abilene 2000     6   156 13910000 66900     780     6.6 2000.
## 7 Abilene 2000     7   152 12635000 73500     742     6.2 2000.
## 8 Abilene 2000     8   131 10710000 75000     765     6.4 2001.
## 9 Abilene 2000     9   104 7615000 64500     771     6.5 2001.
## 10 Abilene 2000    10   101 7040000 59300     764     6.6 2001.
## # ... with 1,646 more rows
```

```
txhousing[,c(1:4)]
```

```
## # A tibble: 8,602 x 4
##   city    year month sales
##   <chr>  <int> <int> <dbl>
## 1 Abilene 2000     1    72
## 2 Abilene 2000     2    98
## 3 Abilene 2000     3   130
## 4 Abilene 2000     4    98
## 5 Abilene 2000     5   141
## 6 Abilene 2000     6   156
## 7 Abilene 2000     7   152
```

```
## 8 Abilene 2000 8 131
## 9 Abilene 2000 9 104
## 10 Abilene 2000 10 101
## # ... with 8,592 more rows
```

```
txhousing[txhousing$city=="San Antonio",c(1:6,8)]
```

```
## # A tibble: 187 x 7
##   city      year month sales    volume median inventory
##   <chr>    <int> <int> <dbl>    <dbl>  <dbl>    <dbl>
## 1 San Antonio 2000 1 820 98974924 90900 4.7
## 2 San Antonio 2000 2 1075 120851076 86000 4.7
## 3 San Antonio 2000 3 1433 167748201 87000 4.9
## 4 San Antonio 2000 4 1263 145280248 90200 5
## 5 San Antonio 2000 5 1574 183281564 91200 5
## 6 San Antonio 2000 6 1666 210779154 100100 5
## 7 San Antonio 2000 7 1508 185816640 100500 4.9
## 8 San Antonio 2000 8 1626 195515195 93400 5.2
## 9 San Antonio 2000 9 1300 156643797 94800 5.2
## 10 San Antonio 2000 10 1192 141630200 93500 5.2
## # ... with 177 more rows
```

```
AbilineData <- txhousing[txhousing$city == "Abilene",]
```

This could be an entire lecture by itself!!! It is important to know how R's indexing works, but in the year 2019 there is no need to be using base R command to index. We will talk more about the tidyverse tomorrow, but the following code does the exact same indexing as the base R code above, but is much more human readable.

Tidyverse

```
txhousing %>%
  filter(year < 2003)
```

```
## # A tibble: 1,656 x 9
##   city      year month sales    volume median listings inventory date
##   <chr>    <int> <int> <dbl>    <dbl>  <dbl>    <dbl>    <dbl>
## 1 Abilene 2000 1 72 5380000 71400 701 6.3 2000
## 2 Abilene 2000 2 98 6505000 58700 746 6.6 2000.
## 3 Abilene 2000 3 130 9285000 58100 784 6.8 2000.
## 4 Abilene 2000 4 98 9730000 68600 785 6.9 2000.
## 5 Abilene 2000 5 141 10590000 67300 794 6.8 2000.
## 6 Abilene 2000 6 156 13910000 66900 780 6.6 2000.
## 7 Abilene 2000 7 152 12635000 73500 742 6.2 2000.
## 8 Abilene 2000 8 131 10710000 75000 765 6.4 2001.
## 9 Abilene 2000 9 104 7615000 64500 771 6.5 2001.
## 10 Abilene 2000 10 101 7040000 59300 764 6.6 2001.
## # ... with 1,646 more rows
```

```
txhousing %>%
  select(city:volume)
```

```
## # A tibble: 8,602 x 5
##   city      year month sales    volume
##   <chr>    <int> <int> <dbl>    <dbl>
```



```
## 1 Abilene 2000 1 72 5380000
## 2 Abilene 2000 2 98 6505000
## 3 Abilene 2000 3 130 9285000
## 4 Abilene 2000 4 98 9730000
## 5 Abilene 2000 5 141 10590000
## 6 Abilene 2000 6 156 13910000
## 7 Abilene 2000 7 152 12635000
## 8 Abilene 2000 8 131 10710000
## 9 Abilene 2000 9 104 7615000
## 10 Abilene 2000 10 101 7040000
## # ... with 8,592 more rows
```

```
txhousing %>%
  select(1:6, inventory) %>%
  filter(city == "San Antonio")
```

```
## # A tibble: 187 x 7
##   city      year month sales    volume median inventory
##   <chr>    <int> <int> <dbl>    <dbl>   <dbl>    <dbl>
## 1 San Antonio 2000 1 820 98974924 90900 4.7
## 2 San Antonio 2000 2 1075 120851076 86000 4.7
## 3 San Antonio 2000 3 1433 167748201 87000 4.9
## 4 San Antonio 2000 4 1263 145280248 90200 5
## 5 San Antonio 2000 5 1574 183281564 91200 5
## 6 San Antonio 2000 6 1666 210779154 100100 5
## 7 San Antonio 2000 7 1508 185816640 100500 4.9
## 8 San Antonio 2000 8 1626 195515195 93400 5.2
## 9 San Antonio 2000 9 1300 156643797 94800 5.2
## 10 San Antonio 2000 10 1192 141630200 93500 5.2
## # ... with 177 more rows
```

```
AbilineData <- txhousing %>%
  filter(city == "Abilene")
```

As your code gets longer, the tidyverse becomes more readable. It is also more helpful for exploring data sets.

Saving and Importing

Finally, if we want to Import or Save other data, we can do that via the Console.

The Working Directory

Most of the work we have done this far is data that we do not want to save. Most of the work you will do after this workshop, you will want to save.

R works by pointing at a folder or directory on your computer. To see where R is pointing now, run the following code

```
getwd()
```

```
## [1] "/Users/davidjohnbaker/Desktop/projects/flash_R"
```

Whatever you **do** in your R session will happen here unless you tell it to otherwise. If you do not want R pointing in this location in your computer, you need to set your working directory elsewhere. To do this, use the `setwd()` command. This is also a good chance to use RStudio's auto complete feature.

```
setwd()
```

Open a double quotation in the function then press TAB. This will allow you to navigate your computer. Going deeper into your directory structure can be done by just following the auto complete. Going higher in the directory requires you to type “which will allow you to look up a level. Set your working directory to theoutput” directory.

The console should now read that it is pointed to the output directory.

You can write a dataset to your working directory with the `write_csv()` command.

```
write.csv(x = AbilineData, file = "MyData.csv")
```

Importing Data

Data is imported using the same logic. You can use the `read.csv()` function to read in a csv file. At first, it might be easier to use the Import Dataset function in RStudio (Top right pane).

Session 3 - tidyverse (1 Hour)



Lesson Goals

- Be able to explain tidy data
- Explain the five tidyverse verbs
- Perform basic indexing
- Import and Export data from R

tidyverse + tidydata

One of the most important concepts data science and R is the idea of tidydata.

The idea behind tidy data is that...

1. Each variable forms a column
2. Each observation forms a row.
3. Each type of observation unit forms a table.

If your data is in this format, then you can do almost anything with the tidyverse.

In order to use the tidyverse, you first need to install it.

```
# install.packages("tidyverse") # Only need to do this once!
library(tidyverse)
```

Five Verbs

The five tidyverse verbs come from the dplyr package. More information on this package can be found here along with these descriptions.

- `mutate()` adds new variables that are functions of existing variables
- `select()` picks variables based on their names.
- `filter()` picks cases based on their values.
- `summarise()` reduces multiple values down to a single summary.
- `arrange()` changes the ordering of the rows.

We can think the verbs as happening in the logical order you would want to grab them. Each of the verbs is also going to be connected to one another with the pipe operator. The idea behind the pipe or “`%>%`” is that the output of the last line is the first argument of the new function.

For example, if we wanted to make a small table that only had data from El Paso from 2011, then only get the first and fifth columns we would run the following code:

```
str(txhousing)

## Classes 'tbl_df', 'tbl' and 'data.frame':   8602 obs. of  9 variables:
##  $ city      : chr  "Abilene" "Abilene" "Abilene" "Abilene" ...
##  $ year      : int   2000 2000 2000 2000 2000 2000 2000 2000 2000 2000 ...
##  $ month     : int    1  2  3  4  5  6  7  8  9 10 ...
##  $ sales     : num   72  98 130  98 141 156 152 131 104 101 ...
##  $ volume    : num  5380000 6505000 9285000 9730000 10590000 ...
##  $ median    : num   71400 58700 58100 68600 67300 66900 73500 75000 64500 59300 ...
##  $ listings  : num   701  746  784  785  794  780  742  765  771  764 ...
##  $ inventory : num    6.3  6.6  6.8  6.9  6.8  6.6  6.2  6.4  6.5  6.6 ...
##  $ date      : num   2000 2000 2000 2000 2000 ...

txhousing_only_el_paso <- txhousing[txhousing$city == "El Paso",]
iris_only_only_el_paso_2005_2011 <- txhousing_only_el_paso[txhousing_only_el_paso$year >= 2011,]
iris_only_only_el_paso_2005_2011[,c(1,5)]

## # A tibble: 55 x 2
##   city      volume
##   <chr>      <dbl>
## 1 El Paso  66136913
## 2 El Paso  44840808
## 3 El Paso  63884923
## 4 El Paso  74429226
## 5 El Paso  61624856
## 6 El Paso  71212091
## 7 El Paso  72366107
## 8 El Paso  86547783
## 9 El Paso  64083406
## 10 El Paso 67015215
## # ... with 45 more rows
```

Which is a bit verbose.

In order to do this with the tidyverse, you would start with the dataset, then run two verbs over it, connected with the pipe.

```
library(tidyverse)
iris_tibble <- as_tibble(iris)

txhousing %>%
  filter(city == "El Paso") %>%
  filter(year >= 2011) %>%
  select(1,5)
```

```
## # A tibble: 55 x 2
##   city      volume
##   <chr>      <dbl>
## 1 El Paso 66136913
## 2 El Paso 44840808
## 3 El Paso 63884923
## 4 El Paso 74429226
## 5 El Paso 61624856
## 6 El Paso 71212091
## 7 El Paso 72366107
## 8 El Paso 86547783
## 9 El Paso 64083406
## 10 El Paso 67015215
## # ... with 45 more rows
```

Both create the same output, but one is much easier to read.

We will now explore a dataset using the five verbs in the dplyr package. You use each of the five verbs as you would in English to think about how you want to manipulate your data.

The key to using the tidyverse is the %>% operator (the pipe operator). It works by taking output from what is before it and piping it to the next command.

Economics Data

The dataset here comes from housing sales data in Texas provided by the TAMU real estate centre.

Variable	Description
city	Name of MLS area
year,month,date	Date
sales	Number of sales
volume	Total value of sales
median	Median sale price
listings	Total active listings
inventory	“Months inventory” : amount of time it would take to sell all current listings at current pace of sales.

Select

The select command works by “Selecting” the columns you wish to work with. It can either take the index of the column using numbers, or the text. There are other options like asking for columns that start or end

with certain text.

```
txhousing %>%  
  select(1,2)
```

```
## # A tibble: 8,602 x 2  
##   city      year  
##   <chr>    <int>  
## 1 Abilene  2000  
## 2 Abilene  2000  
## 3 Abilene  2000  
## 4 Abilene  2000  
## 5 Abilene  2000  
## 6 Abilene  2000  
## 7 Abilene  2000  
## 8 Abilene  2000  
## 9 Abilene  2000  
## 10 Abilene 2000  
## # ... with 8,592 more rows
```

```
txhousing %>%  
  select(1:3)
```

```
## # A tibble: 8,602 x 3  
##   city      year month  
##   <chr>    <int> <int>  
## 1 Abilene  2000     1  
## 2 Abilene  2000     2  
## 3 Abilene  2000     3  
## 4 Abilene  2000     4  
## 5 Abilene  2000     5  
## 6 Abilene  2000     6  
## 7 Abilene  2000     7  
## 8 Abilene  2000     8  
## 9 Abilene  2000     9  
## 10 Abilene 2000    10  
## # ... with 8,592 more rows
```

```
txhousing %>%  
  select(city, sales:median)
```

```
## # A tibble: 8,602 x 4  
##   city      sales  volume median  
##   <chr>    <dbl>    <dbl> <dbl>  
## 1 Abilene     72 5380000  71400  
## 2 Abilene     98 6505000  58700  
## 3 Abilene    130 9285000  58100  
## 4 Abilene     98 9730000  68600  
## 5 Abilene    141 10590000 67300  
## 6 Abilene    156 13910000 66900  
## 7 Abilene    152 12635000 73500  
## 8 Abilene    131 10710000 75000  
## 9 Abilene    104  7615000 64500  
## 10 Abilene    101  7040000 59300  
## # ... with 8,592 more rows
```

```
txhousing %>%
  select(starts_with("ci"))
```

```
## # A tibble: 8,602 x 1
##   city
##   <chr>
## 1 Abilene
## 2 Abilene
## 3 Abilene
## 4 Abilene
## 5 Abilene
## 6 Abilene
## 7 Abilene
## 8 Abilene
## 9 Abilene
## 10 Abilene
## # ... with 8,592 more rows
```

```
txhousing %>%
  select(-city)
```

```
## # A tibble: 8,602 x 8
##   year month sales    volume median listings inventory  date
##   <int> <int> <dbl>    <dbl>    <dbl>    <dbl>    <dbl> <dbl>
## 1 2000     1    72 5380000    71400     701      6.3 2000
## 2 2000     2    98 6505000    58700     746      6.6 2000.
## 3 2000     3   130 9285000    58100     784      6.8 2000.
## 4 2000     4    98 9730000    68600     785      6.9 2000.
## 5 2000     5   141 10590000   67300     794      6.8 2000.
## 6 2000     6   156 13910000   66900     780      6.6 2000.
## 7 2000     7   152 12635000   73500     742      6.2 2000.
## 8 2000     8   131 10710000   75000     765      6.4 2001.
## 9 2000     9   104 7615000    64500     771      6.5 2001.
## 10 2000    10   101 7040000    59300     764      6.6 2001.
## # ... with 8,592 more rows
```

Filter

Once we have the columns we want to work with, we can then pick the rows that are of interest. We do this with the filter function. Here when asking for matches of character strings, you need to use the `==`. R will remind you if you forget. The filter command can be combined with the logical operators. Remember this includes negation operators.

```
txhousing %>%
  filter(city == "El Paso")
```

```
## # A tibble: 187 x 9
##   city    year month sales    volume median listings inventory  date
##   <chr>    <int> <int> <dbl>    <dbl>    <dbl>    <dbl>    <dbl> <dbl>
## 1 El Paso 2000     1   306 31525000  82100    2512      5.8 2000
## 2 El Paso 2000     2   346 32300000  76600    2572      5.9 2000.
## 3 El Paso 2000     3   492 47505000  77100    2549      5.8 2000.
## 4 El Paso 2000     4   382 37915000  79400    2525      5.9 2000.
## 5 El Paso 2000     5   459 43335000  80100    2552      5.9 2000.
## 6 El Paso 2000     6   486 47880000  83200     NA      NA    2000.
```

```
## 7 El Paso 2000 7 422 42925000 82600 2685 6.4 2000.
## 8 El Paso 2000 8 538 53800000 81700 3396 8 2001.
## 9 El Paso 2000 9 382 36775000 78300 2661 6.3 2001.
## 10 El Paso 2000 10 392 40535000 81900 2704 6.5 2001.
## # ... with 177 more rows
```

```
txhousing %>%
  filter(city == "El Paso" | city == "San Antonio")
```

```
## # A tibble: 374 x 9
##   city      year month sales    volume median listings inventory date
##   <chr>    <int> <int> <dbl>    <dbl> <dbl>    <dbl>    <dbl> <dbl>
## 1 El Paso 2000     1  306 31525000 82100    2512     5.8 2000
## 2 El Paso 2000     2  346 32300000 76600    2572     5.9 2000.
## 3 El Paso 2000     3  492 47505000 77100    2549     5.8 2000.
## 4 El Paso 2000     4  382 37915000 79400    2525     5.9 2000.
## 5 El Paso 2000     5  459 43335000 80100    2552     5.9 2000.
## 6 El Paso 2000     6  486 47880000 83200     NA     NA 2000.
## 7 El Paso 2000     7  422 42925000 82600    2685     6.4 2000.
## 8 El Paso 2000     8  538 53800000 81700    3396     8 2001.
## 9 El Paso 2000     9  382 36775000 78300    2661     6.3 2001.
## 10 El Paso 2000    10  392 40535000 81900    2704     6.5 2001.
## # ... with 364 more rows
```

```
txhousing %>%
  filter(city == "El Paso" | city == "San Antonio") %>%
  filter(year >= 2004)
```

```
## # A tibble: 278 x 9
##   city      year month sales    volume median listings inventory date
##   <chr>    <int> <int> <dbl>    <dbl> <dbl>    <dbl>    <dbl> <dbl>
## 1 El Paso 2004     1  435 48330000 93500    3028     5.8 2004
## 2 El Paso 2004     2  441 48215000 89600    3162     6 2004.
## 3 El Paso 2004     3  551 60105000 89000    3288     6.2 2004.
## 4 El Paso 2004     4  579 64980000 89900    3320     6.2 2004.
## 5 El Paso 2004     5  576 68890000 97500    3271     6.1 2004.
## 6 El Paso 2004     6  586 72925000 97200     NA     NA 2004.
## 7 El Paso 2004     7  619 77745000 99900     NA     NA 2004.
## 8 El Paso 2004     8  462 54045000 95400     NA     NA 2005.
## 9 El Paso 2004     9  294 29910000 85800     NA     NA 2005.
## 10 El Paso 2004    10  484 56625000 94800     NA     NA 2005.
## # ... with 268 more rows
```

```
txhousing %>%
  filter(city == "El Paso" | city == "San Antonio") %>%
  filter(year >= 2004) %>%
  filter(month != 1)
```

```
## # A tibble: 254 x 9
##   city      year month sales    volume median listings inventory date
##   <chr>    <int> <int> <dbl>    <dbl> <dbl>    <dbl>    <dbl> <dbl>
## 1 El Paso 2004     2  441 48215000 89600    3162     6 2004.
## 2 El Paso 2004     3  551 60105000 89000    3288     6.2 2004.
## 3 El Paso 2004     4  579 64980000 89900    3320     6.2 2004.
## 4 El Paso 2004     5  576 68890000 97500    3271     6.1 2004.
## 5 El Paso 2004     6  586 72925000 97200     NA     NA 2004.
```

```
## 6 El Paso 2004 7 619 77745000 99900 NA NA 2004.
## 7 El Paso 2004 8 462 54045000 95400 NA NA 2005.
## 8 El Paso 2004 9 294 29910000 85800 NA NA 2005.
## 9 El Paso 2004 10 484 56625000 94800 NA NA 2005.
## 10 El Paso 2004 11 470 55690000 96200 NA NA 2005.
## # ... with 244 more rows
```

Mutate

The mutate command will create new variables.

```
txhousing %>%
  mutate(zSales = scale(sales))
```

```
## # A tibble: 8,602 x 10
##   city   year month sales volume median listings inventory date
##   <chr> <int> <int> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 Abil~ 2000 1 72 5.38e6 71400 701 6.3 2000
## 2 Abil~ 2000 2 98 6.50e6 58700 746 6.6 2000.
## 3 Abil~ 2000 3 130 9.28e6 58100 784 6.8 2000.
## 4 Abil~ 2000 4 98 9.73e6 68600 785 6.9 2000.
## 5 Abil~ 2000 5 141 1.06e7 67300 794 6.8 2000.
## 6 Abil~ 2000 6 156 1.39e7 66900 780 6.6 2000.
## 7 Abil~ 2000 7 152 1.26e7 73500 742 6.2 2000.
## 8 Abil~ 2000 8 131 1.07e7 75000 765 6.4 2001.
## 9 Abil~ 2000 9 104 7.62e6 64500 771 6.5 2001.
## 10 Abil~ 2000 10 101 7.04e6 59300 764 6.6 2001.
## # ... with 8,592 more rows, and 1 more variable: zSales[,1] <dbl>
```

```
txhousing %>%
  filter(city == "El Paso" | city == "San Antonio") %>%
  filter(year >= 2004) %>%
  filter(month != 1) %>%
  mutate(zScale = scale(sales))
```

```
## # A tibble: 254 x 10
##   city   year month sales volume median listings inventory date
##   <chr> <int> <int> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 El P~ 2004 2 441 4.82e7 89600 3162 6 2004.
## 2 El P~ 2004 3 551 6.01e7 89000 3288 6.2 2004.
## 3 El P~ 2004 4 579 6.50e7 89900 3320 6.2 2004.
## 4 El P~ 2004 5 576 6.89e7 97500 3271 6.1 2004.
## 5 El P~ 2004 6 586 7.29e7 97200 NA NA 2004.
## 6 El P~ 2004 7 619 7.77e7 99900 NA NA 2004.
## 7 El P~ 2004 8 462 5.40e7 95400 NA NA 2005.
## 8 El P~ 2004 9 294 2.99e7 85800 NA NA 2005.
## 9 El P~ 2004 10 484 5.66e7 94800 NA NA 2005.
## 10 El P~ 2004 11 470 5.57e7 96200 NA NA 2005.
## # ... with 244 more rows, and 1 more variable: zScale[,1] <dbl>
```

Arrange

Arrange will sort our data.


```
txhousing %>%
  filter(city == "El Paso" | city == "San Antonio") %>%
  filter(year >= 2004) %>%
  filter(month != 1) %>%
  mutate(zScale = scale(sales)) %>%
  arrange(sales)

## # A tibble: 254 x 10
##   city   year month sales volume median listings inventory date
##   <chr> <int> <int> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 El P~ 2011     2  287 4.48e7 134500    3023     6.5 2011.
## 2 El P~ 2008     4  292 4.45e7 126700    4840    10.9 2008.
## 3 El P~ 2004     9  294 2.99e7  85800     NA     NA 2005.
## 4 El P~ 2009     2  326 4.97e7 130900    4530    11.4 2009.
## 5 El P~ 2008     2  328 5.34e7 133800    4374     9 2008.
## 6 El P~ 2013     2  350 5.42e7 139000    3425     7.3 2013.
## 7 El P~ 2005     9  356 4.23e7 111300     NA     NA 2006.
## 8 El P~ 2007    12  362 5.56e7 133500    4625     8.9 2008.
## 9 El P~ 2008    11  362 5.65e7 132800    4773    12 2009.
## 10 El P~ 2008    12  363 5.80e7 136300    4454    11.2 2009.
## # ... with 244 more rows, and 1 more variable: zScale[,1] <dbl>

txhousing %>%
  filter(city == "El Paso" | city == "San Antonio") %>%
  filter(year >= 2004) %>%
  filter(month != 1) %>%
  mutate(zScale = scale(sales)) %>%
  arrange(sales, -year)

## # A tibble: 254 x 10
##   city   year month sales volume median listings inventory date
##   <chr> <int> <int> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 El P~ 2011     2  287 4.48e7 134500    3023     6.5 2011.
## 2 El P~ 2008     4  292 4.45e7 126700    4840    10.9 2008.
## 3 El P~ 2004     9  294 2.99e7  85800     NA     NA 2005.
## 4 El P~ 2009     2  326 4.97e7 130900    4530    11.4 2009.
## 5 El P~ 2008     2  328 5.34e7 133800    4374     9 2008.
## 6 El P~ 2013     2  350 5.42e7 139000    3425     7.3 2013.
## 7 El P~ 2005     9  356 4.23e7 111300     NA     NA 2006.
## 8 El P~ 2008    11  362 5.65e7 132800    4773    12 2009.
## 9 El P~ 2007    12  362 5.56e7 133500    4625     8.9 2008.
## 10 El P~ 2010     2  363 5.16e7 126300    3321     7.4 2010.
## # ... with 244 more rows, and 1 more variable: zScale[,1] <dbl>
```

Group By and Sumemrise

Often we also want to perform the same type of calculation on a group in our dataset. For this we need to group our data, then use the summarize command. We can also use the `n()` function to count the number of observations in each group.

```
txhousing %>%
  filter(city == "El Paso" | city == "San Antonio") %>%
  filter(year >= 2004) %>%
  group_by(year) %>%
```

```

summarise(mean = mean(sales))

## # A tibble: 12 x 2
##   year mean
##   <int> <dbl>
## 1  2004 1102.
## 2  2005 1224.
## 3  2006 1381.
## 4  2007 1257.
## 5  2008 1006.
## 6  2009 1004.
## 7  2010 1000.
## 8  2011  980.
## 9  2012 1090.
## 10 2013 1246.
## 11 2014 1323.
## 12 2015 1461.

txhousing %>%
  filter(city == "El Paso" | city == "San Antonio") %>%
  filter(year >= 2004) %>%
  group_by(year) %>%
  summarise(mean = mean(sales), n = n())

## # A tibble: 12 x 3
##   year mean    n
##   <int> <dbl> <int>
## 1  2004 1102.    24
## 2  2005 1224.    24
## 3  2006 1381.    24
## 4  2007 1257.    24
## 5  2008 1006.    24
## 6  2009 1004.    24
## 7  2010 1000.    24
## 8  2011  980.    24
## 9  2012 1090.    24
## 10 2013 1246.    24
## 11 2014 1323.    24
## 12 2015 1461.    14

```

Work Time

We will now explore the dataset using some guided questions.

In a new script, create following:

- Create a table with only the the first four counties in the dataset.
- Next, run the same command and run that only using one argument that adds in counties that have the work “County” in the title
- Create any new table using a single logical operator
- Create a table with a two logical operators
- Create a table that has no observations from either Paris or Waco.
- Create a new variable based on two other variables
- Find the month with the highest average scales in Tyler county for the year 2015

- Create a table with data from Austin and Galveston, using only the last three years of the dataset. Group the sales by county and then calculate z scores for each county.
- Save your new table to a csv file